

STRATEGY FOR MONITORING PRODUCTION SERVERS

Contents	1
1 Introduction	2
2 Monitoring Tools & Technologies	2
3 Implementation	2
4 Alerts	2

1 Introduction

When faced with the task of monitoring production Ubuntu instances, three important considerations come to mind for me:

1. Choice of tools & technologies to use.
2. How to implement the chosen technologies as a running system
3. Setting up proper channels to alert the relevant administrators of any possible issues.

Below, I will briefly discuss each of these in more detail.

2 Monitoring Tools & Technologies

I have created and configured metrics & monitoring previously, using a *Prometheus-Grafana* stack (and in one case an additional *blackbox_exporter* component for monitoring legacy web and database servers), and this has worked extremely well.

3 Implementation

My implementation of choice for the monitoring stack would be

- Deploying the monitoring stack in Kubernetes.
- Using Helm charts for deploying the different components.

Alternatively, I'd been considering the creation of a Docker cluster (docker compose) for the monitoring stack if Kubernetes or equivalent was not available to me.

4 Alerts

The third consideration would be to set up instant alerting with high visibility.

In my experience email alerts, while useful for detail analysis & troubleshooting, lacks immediacy: People tend to ignore emails or they are simply too occupied with other work to notice an email alert in reasonable time.

I would therefore prefer to set up an instant alert channel using tools such as *MS Teams*, *Opsgenie*, etc. together with email alerting.

I think it is also very important that the metrics set up in Prometheus & Grafana, should be relevant, accurate, and as easy as possible for a viewer to understand; hand-in-hand with this, I would take care that alert messages are accurate and easy to understand, rather than vague and/or misleading.